

Problem Statement

- Scaler is an online tech-versity offering intensive computer science & Data Science courses through live classes delivered by tech leaders and subject matter experts.
- The meticulously structured program enhances the skills of software professionals by offering a modern curriculum with exposure to the latest technologies. It is a product by InterviewBit.
- You are working as a data scientist with the analytics vertical of Scaler, focused on profiling the best companies and job positions to work for from the Scaler database.
- You are provided with the information for a segment of learners and tasked to cluster them on the basis of their job profile, company, and other features. Ideally, these clusters should have similar characteristics.

Data Dictionary:

- 'Unnamed 0'- Index of the dataset
- Email_hash- Anonymised Personal Identifiable Information (PII)
- Company_hash- Current employer of the learner
- orgyear- Employment start date
- CTC- Current CTC
- Job_position- Job profile in the company
- CTC_updated_year: Year in which CTC got updated (Yearly increments, Promotions)

Concept Used:

- Manual Clustering
- Unsupervised Clustering - K- means, Hierarchical Clustering

```
In [127]: import re
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (12,8)
```

```
In [128]: import warnings
warnings.filterwarnings("ignore")
```

```
In [129]: df = pd.read_csv("scaler_clustering.csv",index_col=0)
```

```
In [130]: df.sample(10)
```

```
Out[130]:
```

| | company_hash | email_hash | orgyear | ctc | job_position | ctc_updated_year |
|--------|-----------------------|--|---------|---------|--------------------|------------------|
| 17621 | bvi ogenfvqt | e93abc6cafbdb171f08953540ecf510f10dd3c29698fe2d... | 2015.0 | 200000 | Frontend Engineer | 2021.0 |
| 43264 | qfo | 95359fcf297402a0fd09a5d467e90647494e1820fb4091... | 2018.0 | 600000 | NaN | 2021.0 |
| 6416 | zgzt | fb69e1bf6d85b39e4759ad3db8a1a55c1175c240108cca... | 2016.0 | 450000 | Devops Engineer | 2020.0 |
| 35431 | fytrotjt ntwyzgrgsj | 216b7f3bd41a215b0ff15baef9a2253a8eba2fd0127b7... | 2018.0 | 200000 | NaN | 2021.0 |
| 109059 | xzegojo | 630b0d4ce7833b3a0f4985f36ea19b76c483523be204b6... | 2020.0 | 525000 | FullStack Engineer | 2021.0 |
| 2417 | qxv vacxogqj | adf6018a5bdfcd819beb86808e9c3ed2ea954a543f7dbf... | 2020.0 | 700000 | NaN | 2021.0 |
| 175594 | sggssrt | 8e4b39577f3b328db8ef87cbc841a9fa18be0983157416... | 2018.0 | 1950000 | Frontend Engineer | 2020.0 |
| 47541 | hztburgjta | b4a2b543479e569ccb4591e4490f7685b0856540c08094... | 2018.0 | 24000 | NaN | 2020.0 |
| 84285 | ovu | 08a1ffc2306b7b84edb7081c030c34df39858269cdcd2a... | 2015.0 | 930000 | Frontend Engineer | 2018.0 |
| 134205 | gzbgmxrtsrgmvr rxbxta | 491c9b3c8df401e916538f4d9d39c8a3fee1a39d7db834... | 2012.0 | 1700000 | FullStack Engineer | 2018.0 |

```
In [ ]:
```

```
In [131]: df.shape
```

```
Out[131]: (205843, 6)
```

```
In [ ]: # 205843 learners data
```

```
In [132]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 205843 entries, 0 to 206922
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   company_hash    205799 non-null   object 
 1   email_hash      205843 non-null   object 
 2   orgyear         205757 non-null   float64
 3   ctc             205843 non-null   int64  
 4   job_position    153281 non-null   object 
 5   ctc_updated_year 205843 non-null   float64
dtypes: float64(2), int64(1), object(3)
memory usage: 11.0+ MB
```

```
In [133]: df.isna().sum()
```

```
Out[133]: company_hash      44
email_hash          0
orgyear            86
ctc                0
job_position       52562
ctc_updated_year   0
dtype: int64
```

```
In [134]: (df.isna().sum() / len(df))*100
```

```
Out[134]: company_hash      0.021376
email_hash        0.000000
orgyear          0.041779
ctc              0.000000
job_position     25.534995
ctc_updated_year  0.000000
dtype: float64
```

```
In [ ]:
```

```
In [135]: df.describe()
```

```
Out[135]:      orgyear       ctc  ctc_updated_year
count  205757.000000  2.058430e+05  205843.000000
mean   2014.882750  2.271685e+06  2019.628231
std    63.571115  1.180091e+07   1.325104
min    0.000000  2.000000e+00  2015.000000
25%   2013.000000  5.300000e+05  2019.000000
50%   2016.000000  9.500000e+05  2020.000000
75%   2018.000000  1.700000e+06  2021.000000
max   20165.000000  1.000150e+09  2021.000000
```

```
In [136]: # based on above information , noticing some unusual outliers in the data
```

```
In [ ]:
```

```
In [137]: df.describe(include="object")
```

```
Out[137]:      company_hash           email_hash  job_position
count            205799                  205843        153281
unique           37299                  153443        1017
top  nvnv wgzohrnvwj otqoxwto  bbace3cc586400bbc65765bc6a16b77d8913836cf98b7...  Backend Engineer
freq             8337                      10          43554
```

```
In [ ]:
```

```
In [138]: def preprocess_string(string):
    new_string= re.sub('[^A-Za-z ]+', '', string).lower().strip()
    return new_string

mystring='\tAirtel\\\\\\&**() X Labs'
preprocess_string(mystring)
```

```
Out[138]: 'airtel x labs'
```

```
In [139]: df["company_hash"].nunique()
```

```
Out[139]: 37299
```

```
In [140]: df["company_hash"] = df["company_hash"].apply(lambda x: preprocess_string(str(x)))
df["company_hash"].nunique()
```

```
Out[140]: 37208
```

```
In [ ]:
```

```
In [141]: df["job_position"].nunique()
# 1017 unique job positions are there in the dataset
```

```
Out[141]: 1017
```

```
In [142]: df["job_position"] = df["job_position"].apply(lambda x: preprocess_string(str(x)))
df["job_position"].nunique()

# 857 unique job positions are there in the dataset after preprocessing strings
```

```
Out[142]: 857
```

```
In [ ]:
```

```
In [143]: # removing the email_hash
df.drop("email_hash",axis = 1,inplace=True)
```

```
In [ ]:
```

```
In [144]: df.sample(5)
```

```
Out[144]:
```

| | company_hash | orgyear | ctc | job_position | ctc_updated_year |
|--------|---------------------|---------|---------|--------------------|------------------|
| 135202 | nxmlwg ogenfvqt xzw | 2014.0 | 270000 | backend engineer | 2016.0 |
| 8043 | st | 2012.0 | 1320000 | backend engineer | 2019.0 |
| 191058 | vbkqz rvm | 2010.0 | 220000 | fullstack engineer | 2019.0 |
| 190577 | vagmt | 2016.0 | 2200000 | devops engineer | 2019.0 |
| 64059 | obvqnqgz | 2014.0 | 650000 | android engineer | 2019.0 |

```
In [145]: df.duplicated().sum() # 17597 duplicated records
```

```
Out[145]: 17597
```

```
In [ ]:
```

```
In [ ]:
```

```
In [146]: df.isna().sum()
```

```
Out[146]: company_hash      0
orgyear        86
ctc           0
job_position    0
ctc_updated_year  0
dtype: int64
```

```
In [147]: (df["company_hash"] == "").sum()
```

```
Out[147]: 89
```

```
In [148]: (df["company_hash"] == "nan").sum()
```

```
Out[148]: 44
```

```
In [149]: (df["job_position"] == "").sum()
```

```
Out[149]: 9
```

```
In [150]: (df["job_position"] == "nan").sum()
```

```
Out[150]: 52562
```

```
In [151]: # removing the records where company or job_position reocords are not available
```

```
In [152]: df[(df["company_hash"] == "") | (df["job_position"] == "")].sample(10)
```

```
Out[152]:   company_hash  orgyear     ctc  job_position  ctc_updated_year
0       167717    2018.0  1500000  backend engineer    2020.0
1        76907    2021.0   800000        nan    2021.0
2       25333    2019.0  2000000        nan    2021.0
3       202179    2016.0   500000        nan    2017.0
4       84192    2018.0  1400000  backend engineer    2019.0
5      197978    2020.0  1000000        nan    2019.0
6       50414    2020.0   720000        nan    2019.0
7      117571    2010.0  4500000        nan    2019.0
8      127679    2019.0  1400000  backend engineer    2019.0
9       80668    2019.0   850000        nan    2019.0
```

```
In [153]: len(df[(df["company_hash"] == "") | (df["job_position"] == "")])
```

```
Out[153]: 98
```

```
In [154]: # df[((df["company_hash"] != "") & (df["job_position"] != ""))]
```

```
In [155]: df = df[~((df["company_hash"] == "") | (df["job_position"] == ""))]
```

```
Out[155]:
```

| | company_hash | orgyear | ctc | job_position | ctc_updated_year |
|--------|-------------------------|---------|---------|--------------------|------------------|
| 0 | atrxnnt xzaxv | 2016.0 | 1100000 | other | 2020.0 |
| 1 | qtxvzwt xzegwgb rxbxnta | 2018.0 | 449999 | fullstack engineer | 2019.0 |
| 2 | ojzwnvwnxw vx | 2015.0 | 2000000 | backend engineer | 2020.0 |
| 3 | ngpgutaxv | 2017.0 | 700000 | backend engineer | 2019.0 |
| 4 | qxen sqghu | 2017.0 | 1400000 | fullstack engineer | 2019.0 |
| ... | ... | ... | ... | ... | ... |
| 206918 | vuurt xzw | 2008.0 | 220000 | nan | 2019.0 |
| 206919 | husqvawgb | 2017.0 | 500000 | nan | 2020.0 |
| 206920 | vwwgrxnt | 2021.0 | 700000 | nan | 2021.0 |
| 206921 | zgn vuurxwvmrt | 2019.0 | 5100000 | nan | 2019.0 |
| 206922 | bgqsvz onvzrtj | 2014.0 | 1240000 | nan | 2016.0 |

205745 rows × 5 columns

Data Preprocessing

```
In [156]: df["orgyear"].isna().sum()
```

```
Out[156]: 86
```

- imputing Employee Start Year as per the median year as per each company.

```
In [157]: df.groupby("company_hash")["orgyear"].transform("median")
```

```
Out[157]:
```

| | |
|--------|--------|
| 0 | 2014.0 |
| 1 | 2016.0 |
| 2 | 2015.0 |
| 3 | 2016.0 |
| 4 | 2017.0 |
| ... | ... |
| 206918 | 2018.0 |
| 206919 | 2017.0 |
| 206920 | 2016.0 |
| 206921 | 2020.0 |
| 206922 | 2015.0 |

Name: orgyear, Length: 205745, dtype: float64

```
In [158]: df["orgyear"].fillna(df['orgyear'].isnull().sum(), inplace=True)
```

```
In [159]: df["orgyear"].isna().sum()
```

```
Out[159]: 0
```

```
In [160]: df.sample(5)
```

```
Out[160]:
```

| | company_hash | orgyear | ctc | job_position | ctc_updated_year |
|--------|--------------------------|---------|---------|--------------------|------------------|
| 175831 | bxqtrk | 2013.0 | 2500000 | fullstack engineer | 2019.0 |
| 49006 | wxnx | 2018.0 | 1500000 | backend engineer | 2021.0 |
| 21098 | tdr | 2015.0 | 730000 | other | 2020.0 |
| 151312 | nvnv wgzohrnvwj otqcxwto | 2020.0 | 700000 | fullstack engineer | 2020.0 |
| 153058 | vwwtzhnq | 2016.0 | 700000 | nan | 2021.0 |

```
In [ ]:
```

```
In [ ]:
```

Outliers Treatment :

- employment start year

```
In [161]: df["orgyear"].value_counts()
```


```
Out[161]:
```

| | |
|--------|-------|
| 2018.0 | 25240 |
| 2019.0 | 23402 |
| 2017.0 | 23237 |
| 2016.0 | 23038 |
| 2015.0 | 20602 |
| ... | |
| 2107.0 | 1 |
| 1972.0 | 1 |
| 2101.0 | 1 |
| 208.0 | 1 |
| 200.0 | 1 |

Name: orgyear, Length: 78, dtype: int64

```
In [162]: sns.countplot(df["orgyear"])
plt.xticks(rotation = 90)
plt.show()
```

<IPython.core.display.Javascript object>



```
In [163]: # sns.histplot(np.log(df["orgyear"]))
```

```
In [164]: df["orgyear"].quantile(0.001)
```

Out[164]: 1990.0


```
In [165]: df["orgyear"].quantile(0.999)
```

Out[165]: 2023.0

```
In [166]: df["orgyear"] = df["orgyear"].clip(1990,2022)
```

```
In [167]: sns.countplot(df["orgyear"])
plt.xticks(rotation = 90)
plt.show()
```

<IPython.core.display.Javascript object>



```
In [ ]:
```

- **ctc updated_year**

```
In [168]: df["ctc_updated_year"].quantile(0.001)
```


Out[168]: 2015.0

```
In [169]: df["ctc_updated_year"].quantile(0.99)
```

Out[169]: 2021.0

```
In [170]: sns.countplot(df["ctc_updated_year"])
plt.xticks(rotation = 90)
plt.show()
```

<IPython.core.display.Javascript object>



```
In [ ]:
```

- outlier treatment for CTC

```
In [171]: df["ctc"].quantile(0.01)
```

Out[171]: 37000.0

```
In [172]: df["ctc"].quantile(0.999)
```

Out[172]: 200000000.0

```
In [173]: df = df.loc[((df.ctc) > df.ctc.quantile(0.01)) & ((df.ctc) < df.ctc.quantile(0.999))]
```

```
In [174]: df
```

```
Out[174]:
```


| | company_hash | orgyear | ctc | job_position | ctc_updated_year |
|--------|-------------------------|---------|---------|--------------------|------------------|
| 0 | atrgxnnt xzaxv | 2016.0 | 1100000 | other | 2020.0 |
| 1 | qtrxzwt xzegwgb rxbxnta | 2018.0 | 449999 | fullstack engineer | 2019.0 |
| 2 | ojzwnvwnxw vx | 2015.0 | 2000000 | backend engineer | 2020.0 |
| 3 | ngpputaxv | 2017.0 | 700000 | backend engineer | 2019.0 |
| 4 | qxen sqghu | 2017.0 | 1400000 | fullstack engineer | 2019.0 |
| ... | ... | ... | ... | ... | ... |
| 206918 | vuurx zzw | 2008.0 | 220000 | nan | 2019.0 |
| 206919 | husqavawgb | 2017.0 | 500000 | nan | 2020.0 |
| 206920 | vwgrxnt | 2021.0 | 700000 | nan | 2021.0 |
| 206921 | zgn vuurxwvmt | 2019.0 | 5100000 | nan | 2019.0 |
| 206922 | bgqsvz onvzrtj | 2014.0 | 1240000 | nan | 2016.0 |

201625 rows × 5 columns

```
In [175]: sns.distplot(df["ctc"])
```

<IPython.core.display.Javascript object>

```
Out[175]: <AxesSubplot:xlabel='ctc', ylabel='Density'>
```



- replacing string "nan" to np.nan

```
In [176]: df.loc[df['job_position']=='nan', 'job_position']=np.nan  
  
In [177]: df.loc[df["company_hash"]=="nan","company_hash"] = np.nan  
  
In [ ]:  
  
In [ ]:  
  
In [ ]:  
  
In [ ]:  
  
In [270]: # df.company_hash.value_counts(dropna=False)  
  
In [271]: # df.job_position.value_counts(dropna=False)
```

Feature Engineering

Masked company name to "Others" having count less than 5

```
In [180]: df.loc[df.groupby("company_hash")["ctc"].transform("count") < 5,"company_hash"] = "Others"  
  
In [181]: (df["company_hash"] == "Others").sum()  
  
Out[181]: 46434  
  
In [272]: # df.company_hash.value_counts(dropna=False)  
  
In [ ]:  
  
In [183]: df['orgyear'].describe()  
  
Out[183]: count    201625.000000  
mean      2015.104769  
std       4.256063  
min      1990.000000  
25%      2013.000000  
50%      2016.000000  
75%      2018.000000  
max      2022.000000  
Name: orgyear, dtype: float64
```

years of experience = current year - employement start year

```
In [184]: # years of experience  
df["years_of_experience_in_organization"] = 2023 - df["orgyear"]
```


```
In [185]: df.sample(2)
```

```
Out[185]:   company_hash orgyear      ctc job_position ctc_updated_year years_of_experience_in_organization
157723        Others    2019.0  480999       NaN        2018.0                  4.0
85654         xzegojo   2018.0  900000     other        2020.0                  5.0
```

```
In [186]: sns.countplot(df["years_of_experience_in_organization"])
```

```
<IPython.core.display.Javascript object>
```

```
Out[186]: <AxesSubplot:xlabel='years_of_experience_in_organization', ylabel='count'>
```



```
In [ ]:
```

```
In [ ]:
```

```
In [187]: df.duplicated().sum()
```

```
Out[187]: 37683
```

```
In [188]: df.drop_duplicates(inplace=True)
df.shape
```

```
Out[188]: (163942, 6)
```

```
In [189]: df.isna().sum()
```

```
Out[189]: company_hash          42
orgyear                  0
ctc                      0
job_position            36745
ctc_updated_year         0
years_of_experience_in_organization 0
dtype: int64
```

treating records having ctc_updated_year higher than their organization joining year

```
In [190]: # records having ctc_updated_year higher than their organization joining year
(df["ctc_updated_year"] < df["orgyear"]).sum()
```

```
Out[190]: 7181
```

```
In [191]: df.ctc_updated_year = df[["ctc_updated_year", "orgyear"]].max(axis = 1)
```

```
In [192]: (df["ctc_updated_year"] < df["orgyear"]).sum()
```

```
Out[192]: 0
```

```
In [193]: df.sample(2)
```

```
Out[193]:   company_hash  orgyear      ctc  job_position  ctc_updated_year  years_of_experience_in_organization
198803    bgqsvz onvzrtj  2017.0  1600000       NaN        2019.0                 6.0
178348    bjznqvivmgzs  2017.0  1970000       NaN        2017.0                 6.0
```

Filling null values with others -- if not done before

```
In [194]: df['job_position'] = df['job_position'].fillna('Others')
df['company_hash'] = df['company_hash'].fillna('Others')
```

```
In [ ]:
```

```
In [195]: df.isna().sum()
```

```
Out[195]: company_hash          0
orgyear                  0
ctc                      0
job_position            0
ctc_updated_year         0
years_of_experience_in_organization 0
dtype: int64
```

```
In [ ]:
```

```
In [ ]:
```

```
In [196]: df.duplicated().sum()
```

```
Out[196]: 1061
```

```
In [197]: # df.drop_duplicates(inplace=True)
```

```
In [273]: # glacing over data after outlier treatment and preprocessing
```

```
In [198]: df.describe()
```

```
Out[198]:
```

| | orgyear | ctc | ctc_updated_year | years_of_experience_in_organization |
|-------|---------------|--------------|------------------|-------------------------------------|
| count | 163942.000000 | 1.639420e+05 | 163942.000000 | 163942.000000 |
| mean | 2014.772218 | 1.425498e+06 | 2019.595540 | 8.227782 |
| std | 4.402053 | 1.303985e+06 | 1.334962 | 4.402053 |
| min | 1990.000000 | 3.800000e+04 | 2015.000000 | 1.000000 |
| 25% | 2013.000000 | 6.000000e+05 | 2019.000000 | 5.000000 |
| 50% | 2016.000000 | 1.039999e+06 | 2020.000000 | 7.000000 |
| 75% | 2018.000000 | 1.800000e+06 | 2021.000000 | 10.000000 |
| max | 2022.000000 | 1.250000e+07 | 2022.000000 | 33.000000 |

```
In [ ]:
```


```
In [199]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 163942 entries, 0 to 206922
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   company_hash    163942 non-null   object 
 1   orgyear         163942 non-null   float64
 2   ctc             163942 non-null   int64  
 3   job_position    163942 non-null   object 
 4   ctc_updated_year 163942 non-null   float64
 5   years_of_experience_in_organization 163942 non-null   float64
dtypes: float64(3), int64(1), object(2)
memory usage: 8.8+ MB
```

```
In [200]: sns.scatterplot(df.ctc,df.years_of_experience_in_organization)
```

```
<IPython.core.display.Javascript object>
```

```
Out[200]: <AxesSubplot:xlabel='ctc', ylabel='years_of_experience_in_organization'>
```



```
In [ ]:
```

```
In [201]: df.columns
```

```
Out[201]: Index(['company_hash', 'orgyear', 'ctc', 'job_position', 'ctc_updated_year',
   'years_of_experience_in_organization'],
  dtype='object')
```

Manual Clustering based on Company , Job position and Years of experience

Learner's "designation_in_organization"

```
In [ ]:
```

```
In [ ]:
```

In []:

In [203]: GROUPED_CTC

Out[203]:

| years_of_experience_in_organization | job_position | company_hash | count | mean | std | min | 25% | 50% | 75% | max |
|-------------------------------------|--------------------|----------------------------------|-------|--------------|--------------|----------|-----------|-----------|-----------|------------|
| 1.0 | Others | Others | 58.0 | 1.586207e+06 | 2.080212e+06 | 60000.0 | 407500.0 | 750000.0 | 1575000.0 | 10000000.0 |
| | | agzn fgqp xz vzej gqsvzxkvnxgz | 1.0 | 1.600000e+06 | | NaN | 1600000.0 | 1600000.0 | 1600000.0 | 1600000.0 |
| | | atrgxnnt | 1.0 | 1.000000e+06 | | NaN | 1000000.0 | 1000000.0 | 1000000.0 | 1000000.0 |
| | | attr | 1.0 | 1.000000e+06 | | NaN | 1000000.0 | 1000000.0 | 1000000.0 | 1000000.0 |
| | | attr ntwyzzgrgsxto | 2.0 | 1.000000e+06 | 2.828427e+05 | 800000.0 | 900000.0 | 1000000.0 | 1100000.0 | 1200000.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 33.0 | qa engineer | hxxtntaytvrny sqghu | 1.0 | 5.400000e+05 | | NaN | 540000.0 | 540000.0 | 540000.0 | 540000.0 |
| | | tmxd ogenfvqt xzaxv ucen rna | 1.0 | 1.220000e+06 | | NaN | 1220000.0 | 1220000.0 | 1220000.0 | 1220000.0 |
| | | utrvnqg ogrhnxzgo ucnrna | 1.0 | 6.000000e+05 | | NaN | 600000.0 | 600000.0 | 600000.0 | 600000.0 |
| | research engineers | ovbohzs qa xzonxnhnt xzaxv atryx | 1.0 | 1.400000e+06 | | NaN | 1400000.0 | 1400000.0 | 1400000.0 | 1400000.0 |
| | support engineer | Others | 2.0 | 3.700000e+05 | 3.252691e+05 | 140000.0 | 255000.0 | 370000.0 | 485000.0 | 600000.0 |

66191 rows × 8 columns

In []:

```
In [205]: df_GROUPED_CTC_BY_E_P_C
```

```
Out[205]:
```

| | company_hash | orgyear | ctc | job_position | ctc_updated_year | years_of_experience_in_organization | count | mean | std | min | 25% | 50% | 75% | max |
|--------|-------------------------|---------|---------|--------------------|------------------|-------------------------------------|-------|--------------------|--------------|-----------|-----------|-----------|-----------|------------|
| 0 | atrgxntt xzaxv | 2016.0 | 1100000 | other | 2020.0 | | 7.0 | 1.0 1.100000e+06 | NaN | 1100000.0 | 1100000.0 | 1100000.0 | 1100000.0 | 1100000.0 |
| 1 | qtrxzwt xzegwgb rxbxnta | 2018.0 | 449999 | fullstack engineer | 2019.0 | | 5.0 | 7.0 7.742856e+05 | 2.509223e+05 | 449999.0 | 610000.0 | 750000.0 | 900000.0 | 1200000.0 |
| 2 | Others | 2015.0 | 2000000 | backend engineer | 2020.0 | | 8.0 | 440.0 1.269393e+06 | 1.405136e+06 | 41000.0 | 400000.0 | 900000.0 | 1600000.0 | 10000000.0 |
| 3 | ngpputavx | 2017.0 | 700000 | backend engineer | 2019.0 | | 6.0 | 7.0 1.158571e+06 | 4.047810e+05 | 700000.0 | 825000.0 | 1200000.0 | 1405000.0 | 1750000.0 |
| 4 | qxen sqghu | 2017.0 | 1400000 | fullstack engineer | 2019.0 | | 6.0 | 1.0 1.400000e+06 | NaN | 1400000.0 | 1400000.0 | 1400000.0 | 1400000.0 | 1400000.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 163937 | vuur tzw | 2008.0 | 220000 | Others | 2019.0 | | 15.0 | 1.0 2.200000e+05 | NaN | 220000.0 | 220000.0 | 220000.0 | 220000.0 | 220000.0 |
| 163938 | husqvawgb | 2017.0 | 500000 | Others | 2020.0 | | 6.0 | 3.0 1.150000e+06 | 5.634714e+05 | 500000.0 | 975000.0 | 1450000.0 | 1475000.0 | 1500000.0 |
| 163939 | vwwgrxt | 2021.0 | 700000 | Others | 2021.0 | | 2.0 | 3.0 6.666667e+05 | 3.511885e+05 | 300000.0 | 500000.0 | 700000.0 | 850000.0 | 1000000.0 |
| 163940 | zgn vuurxwvmt | 2019.0 | 5100000 | Others | 2019.0 | | 4.0 | 118.0 1.412015e+06 | 1.715935e+06 | 45000.0 | 400000.0 | 735000.0 | 1725250.0 | 10000000.0 |
| 163941 | bgqsyz onvzrtj | 2014.0 | 1240000 | Others | 2016.0 | | 9.0 | 9.0 1.693333e+06 | 3.484250e+05 | 1200000.0 | 1500000.0 | 1700000.0 | 1900000.0 | 2200000.0 |

163942 rows × 14 columns

```
In [206]: def classification(x, ctc_50, ctc_75):  
    if x < ctc_50:  
        return 3  
    elif x >= ctc_50 and x <= ctc_75:  
        return 2  
    elif x >= ctc_75:  
        return 1
```

whichever learner has ctc compared to their years of experience , respective company , position

giving designation as 3 when ctc is < 50th percentile in his position ,experience and company

giving designation as 2 when ctc is between 50th and 75th percentile in his position ,experience and company

giving designation as 1 when ctc is > 75th percentile in his position ,experience and company

```
In [207]: df_GROUPED_CTC_BY_E_P_C["designation_in_organization"] = df_GROUPED_CTC_BY_E_P_C.apply(lambda x:classification(x["ctc"],x["50%"],x["75%"]),axis = 1)
```

```
In [208]: df_GROUPED_CTC_BY_E_P_C
```

```
Out[208]:
```

| | company_hash | orgyear | ctc | job_position | ctc_updated_year | years_of_experience_in_organization | count | mean | std | min | 25% | 50% | 75% | max | designation_in_organization | |
|--------|-------------------------------|---------|---------|-----------------------|------------------|-------------------------------------|-------|-------|--------------|--------------|-----------|-----------|-----------|-----------|-----------------------------|---|
| 0 | atrgxnnt xzaxv | 2016.0 | 1100000 | other | 2020.0 | | 7.0 | 1.0 | 1.100000e+06 | NaN | 1100000.0 | 1100000.0 | 1100000.0 | 1100000.0 | 1100000.0 | 2 |
| 1 | qtrxvzwt xzegwgb rbxnta | 2018.0 | 449999 | fullstack engineer | 2019.0 | | 5.0 | 7.0 | 7.742856e+05 | 2.509223e+05 | 449999.0 | 610000.0 | 750000.0 | 900000.0 | 1200000.0 | 3 |
| 2 | Others | 2015.0 | 2000000 | backend engineer | 2020.0 | | 8.0 | 440.0 | 1.269393e+06 | 1.405136e+06 | 41000.0 | 400000.0 | 900000.0 | 1600000.0 | 10000000.0 | 1 |
| 3 | ngpgutaxv | 2017.0 | 700000 | backend engineer | 2019.0 | | 6.0 | 7.0 | 1.158571e+06 | 4.047810e+05 | 700000.0 | 825000.0 | 1200000.0 | 1405000.0 | 1750000.0 | 3 |
| 4 | qxen sqghu | 2017.0 | 1400000 | fullstack engineer | 2019.0 | | 6.0 | 1.0 | 1.400000e+06 | NaN | 1400000.0 | 1400000.0 | 1400000.0 | 1400000.0 | 1400000.0 | 2 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 163937 | vuurt xzw | 2008.0 | 220000 | Others | 2019.0 | | 15.0 | 1.0 | 2.200000e+05 | NaN | 220000.0 | 220000.0 | 220000.0 | 220000.0 | 220000.0 | 2 |
| 163938 | husqyawgb | 2017.0 | 500000 | Others | 2020.0 | | 6.0 | 3.0 | 1.150000e+06 | 5.634714e+05 | 500000.0 | 975000.0 | 1450000.0 | 1475000.0 | 1500000.0 | 3 |
| 163939 | vwwgrxnt | 2021.0 | 700000 | Others | 2021.0 | | 2.0 | 3.0 | 6.666667e+05 | 3.511885e+05 | 300000.0 | 500000.0 | 700000.0 | 850000.0 | 1000000.0 | 2 |
| 163940 | zgn vuurxwmrt | 2019.0 | 5100000 | Others | 2019.0 | | 4.0 | 118.0 | 1.412015e+06 | 1.715935e+06 | 45000.0 | 400000.0 | 735000.0 | 1725250.0 | 10000000.0 | 1 |
| 163941 | bgqsvz onvzrlj | 2014.0 | 1240000 | Others | 2016.0 | | 9.0 | 9.0 | 1.693333e+06 | 3.484250e+05 | 1200000.0 | 1500000.0 | 1700000.0 | 1900000.0 | 2200000.0 | 3 |

163942 rows × 15 columns

```
In [209]: df_GROUPED_CTC_BY_E_P_C.designation_in_organization.value_counts(normalize=True)
```

```
Out[209]:
```

| | |
|---|----------|
| 2 | 0.456393 |
| 3 | 0.331660 |
| 1 | 0.211947 |

Name: designation_in_organization, dtype: float64

```
In [ ]:
```

```
In [ ]:
```

```
In [211]: df_GROUPED_CTC_BY_E_P_C
```

```
Out[211]:
```

| | company_hash | orgyear | ctc | job_position | ctc_updated_year | years_of_experience_in_organization | count | mean | std | min | 25% | 50% | 75% | max | designation_in_organization | |
|--------|------------------------|---------|---------|--------------------|------------------|-------------------------------------|-------|-------|--------------|--------------|-----------|-----------|-----------|-----------|-----------------------------|---|
| 0 | atrgxnnt xzavx | 2016.0 | 1100000 | other | 2020.0 | | 7.0 | 1.0 | 1.100000e+06 | NaN | 1100000.0 | 1100000.0 | 1100000.0 | 1100000.0 | 1100000.0 | 2 |
| 1 | qtrxzwt xzegwgb rxbxta | 2018.0 | 449999 | fullstack engineer | 2019.0 | | 5.0 | 7.0 | 7.742856e+05 | 2.509223e+05 | 449999.0 | 610000.0 | 750000.0 | 900000.0 | 1200000.0 | 3 |
| 2 | Others | 2015.0 | 2000000 | backend engineer | 2020.0 | | 8.0 | 440.0 | 1.269393e+06 | 1.405136e+06 | 41000.0 | 400000.0 | 900000.0 | 1600000.0 | 10000000.0 | 1 |
| 3 | ngpgutaxv | 2017.0 | 700000 | backend engineer | 2019.0 | | 6.0 | 7.0 | 1.158571e+06 | 4.047810e+05 | 700000.0 | 825000.0 | 1200000.0 | 1405000.0 | 1750000.0 | 3 |
| 4 | qxen sqghu | 2017.0 | 1400000 | fullstack engineer | 2019.0 | | 6.0 | 1.0 | 1.400000e+06 | NaN | 1400000.0 | 1400000.0 | 1400000.0 | 1400000.0 | 1400000.0 | 2 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 163937 | vuurt xzw | 2008.0 | 220000 | Others | 2019.0 | | 15.0 | 1.0 | 2.200000e+05 | NaN | 220000.0 | 220000.0 | 220000.0 | 220000.0 | 220000.0 | 2 |
| 163938 | husqvawgb | 2017.0 | 500000 | Others | 2020.0 | | 6.0 | 3.0 | 1.150000e+06 | 5.634714e+05 | 500000.0 | 975000.0 | 1450000.0 | 1475000.0 | 1500000.0 | 3 |
| 163939 | vwwgrxnt | 2021.0 | 700000 | Others | 2021.0 | | 2.0 | 3.0 | 6.666667e+05 | 3.511885e+05 | 300000.0 | 500000.0 | 700000.0 | 850000.0 | 1000000.0 | 2 |
| 163940 | zgn vuurxwmrt | 2019.0 | 5100000 | Others | 2019.0 | | 4.0 | 118.0 | 1.412015e+06 | 1.715935e+06 | 45000.0 | 400000.0 | 735000.0 | 1725250.0 | 10000000.0 | 1 |
| 163941 | bgqsyz onvzrtj | 2014.0 | 1240000 | Others | 2016.0 | | 9.0 | 9.0 | 1.693333e+06 | 3.484250e+05 | 1200000.0 | 1500000.0 | 1700000.0 | 1900000.0 | 2200000.0 | 3 |

163942 rows × 15 columns

```
In [212]: df_GROUPED_CTC_BY_E_P_C.drop(columns=['count',
```

```
'mean',
'std',
'min',
'25%',
'50%',
'75%',
'max'],axis = 1,inplace=True)
```

```
In [213]: df_GROUPED_CTC_BY_E_P_C
```

```
Out[213]:
```

| | company_hash | orgyear | ctc | job_position | ctc_updated_year | years_of_experience_in_organization | designation_in_organization |
|--------|------------------------|---------|---------|--------------------|------------------|-------------------------------------|-----------------------------|
| 0 | atrgxnnt xzavx | 2016.0 | 1100000 | other | 2020.0 | 7.0 | 2 |
| 1 | qtrxzwt xzegwgb rxbxta | 2018.0 | 449999 | fullstack engineer | 2019.0 | 5.0 | 3 |
| 2 | Others | 2015.0 | 2000000 | backend engineer | 2020.0 | 8.0 | 1 |
| 3 | ngpgutaxv | 2017.0 | 700000 | backend engineer | 2019.0 | 6.0 | 3 |
| 4 | qxen sqghu | 2017.0 | 1400000 | fullstack engineer | 2019.0 | 6.0 | 2 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 163937 | vuurt xzw | 2008.0 | 220000 | Others | 2019.0 | 15.0 | 2 |
| 163938 | husqvawgb | 2017.0 | 500000 | Others | 2020.0 | 6.0 | 3 |
| 163939 | vwwgrxnt | 2021.0 | 700000 | Others | 2021.0 | 2.0 | 2 |
| 163940 | zgn vuurxwmrt | 2019.0 | 5100000 | Others | 2019.0 | 4.0 | 1 |
| 163941 | bgqsyz onvzrtj | 2014.0 | 1240000 | Others | 2016.0 | 9.0 | 3 |

163942 rows × 7 columns

```
In [214]: df_GROUPED_CTC_BY_E_P_C.shape
```

```
Out[214]: (163942, 7)
```

Manual Clustering on company and job position

grouping by each job_position and company,

finding which class of job an individual have,

based on his ctc compared to his job_position and respective company.

```
In [215]: GROUPED_C_J=df.groupby(['job_position','company_hash'])['ctc'].describe()  
GROUPED_C_J
```

```
Out[215]:
```

| job_position | company_hash | count | mean | std | min | 25% | 50% | 75% | max |
|-----------------------|---|--------|--------------|--------------|----------|----------|----------|------------|------------|
| Others | Others | 3520.0 | 1.366188e+06 | 1.445330e+06 | 40000.0 | 409999.0 | 900000.0 | 1842499.25 | 12500000.0 |
| | a ntwyzgrgsxto | 6.0 | 1.229167e+06 | 1.401465e+06 | 350000.0 | 518750.0 | 587500.0 | 1162500.0 | 4000000.0 |
| | aaqxctz avnv owxtzwto vzvrjnxwo ucn rma | 1.0 | 5000000e+05 | NaN | 500000.0 | 500000.0 | 500000.0 | 500000.0 | 500000.0 |
| | abwavnv ojontb | 1.0 | 7.000000e+05 | NaN | 700000.0 | 700000.0 | 700000.0 | 700000.0 | 700000.0 |
| | adw ntwyzgrgsj | 69.0 | 8.502319e+05 | 1.036041e+06 | 80000.0 | 380000.0 | 500000.0 | 1000000.0 | 8000000.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| wordpress developer | Others | 1.0 | 6.000000e+05 | NaN | 600000.0 | 600000.0 | 600000.0 | 600000.0 | 600000.0 |
| worker | zgn vuurxwvmrt vwwghzn | 1.0 | 2.000000e+05 | NaN | 200000.0 | 200000.0 | 200000.0 | 200000.0 | 200000.0 |
| x | Others | 1.0 | 4.000000e+05 | NaN | 400000.0 | 400000.0 | 400000.0 | 400000.0 | 400000.0 |
| young professional ii | sgctqzbztn ge xzaxv | 1.0 | 5.000000e+05 | NaN | 500000.0 | 500000.0 | 500000.0 | 500000.0 | 500000.0 |
| zomato | kgbvng | 2.0 | 3.000000e+05 | 2.828427e+05 | 100000.0 | 200000.0 | 300000.0 | 400000.0 | 500000.0 |

25593 rows × 8 columns

```
In [216]: df_GROUPED_C_J=df.merge(GROUPED_C_J, on=['job_position','company_hash'], how='left')
```

```
In [217]: df_GROUPED_C_J.sample(5)
```

```
Out[217]:
```

| | company_hash | orgyear | ctc | job_position | ctc_updated_year | years_of_experience_in_organization | count | mean | std | min | 25% | 50% | 75% | max | |
|--------|----------------------------|---------|---------|------------------|------------------|-------------------------------------|-------|--------|--------------|--------------|----------|----------|-----------|-----------|------------|
| 126677 | zvz | 2019.0 | 3200000 | data scientist | 2019.0 | | 4.0 | 39.0 | 1.211500e+06 | 7.584555e+05 | 46500.0 | 647500.0 | 1000000.0 | 1650000.0 | 3200000.0 |
| 93215 | nnvv wgzohrnvzwyj otqcxwto | 2012.0 | 850000 | ios engineer | 2019.0 | | 11.0 | 19.0 | 6.852632e+05 | 3.520475e+05 | 70000.0 | 425000.0 | 600000.0 | 850000.0 | 1550000.0 |
| 29447 | wvustbxzx | 2013.0 | 910000 | backend engineer | 2021.0 | | 10.0 | 247.0 | 8.295992e+05 | 4.891276e+05 | 40000.0 | 475000.0 | 700000.0 | 1070000.0 | 3000000.0 |
| 41080 | znn avnv otqcxwto | 2019.0 | 700000 | Others | 2019.0 | | 4.0 | 62.0 | 1.142984e+06 | 1.738646e+06 | 300000.0 | 400000.0 | 690000.0 | 1275000.0 | 10000000.0 |
| 76917 | Others | 2013.0 | 1100000 | other | 2021.0 | | 10.0 | 2367.0 | 1.117373e+06 | 1.423744e+06 | 38000.0 | 350000.0 | 700000.0 | 1326000.0 | 12000000.0 |

```
In [218]: # creating classes basis on the salary in their respective company
```

```
In [219]: df_GROUPED_C_J['classss'] = df_GROUPED_C_J.apply(lambda x: classification(x['ctc'],x['50%'],x['75%']),axis=1)
```

```
In [220]: df_GROUPED_C_J.sample(5)
```

```
Out[220]:
```

| | company_hash | orgyear | ctc | job_position | ctc_updated_year | years_of_experience_in_organization | count | mean | std | min | 25% | 50% | 75% | max | classss | |
|--------|----------------------|---------|--------|--------------------|------------------|-------------------------------------|-------|--------|--------------|--------------|----------|-----------|-----------|-----------|------------|---|
| 79071 | ohnytqrvza | 2017.0 | 710000 | data scientist | 2019.0 | | 6.0 | 2.0 | 6.155000e+06 | 7.700393e+06 | 710000.0 | 3432500.0 | 6155000.0 | 8877500.0 | 11600000.0 | 3 |
| 122059 | Others | 2019.0 | 130000 | fullstack engineer | 2020.0 | | 4.0 | 3181.0 | 1.193104e+06 | 1.532888e+06 | 40000.0 | 350000.0 | 775000.0 | 1450000.0 | 12000000.0 | 3 |
| 35726 | wwnho wgbbhzxwvnxgzo | 2013.0 | 400000 | Others | 2020.0 | | 10.0 | 3.0 | 7.333333e+05 | 5.773503e+05 | 400000.0 | 400000.0 | 400000.0 | 900000.0 | 1400000.0 | 2 |
| 68461 | vau | 2015.0 | 819999 | backend engineer | 2019.0 | | 8.0 | 51.0 | 1.124500e+06 | 9.053526e+05 | 105000.0 | 690000.0 | 820000.0 | 1325000.0 | 6000000.0 | 3 |
| 42240 | wgszxkvzn | 2015.0 | 800000 | frontend engineer | 2021.0 | | 8.0 | 105.0 | 7.959143e+05 | 5.380693e+05 | 65000.0 | 450000.0 | 600000.0 | 1000000.0 | 4298000.0 | 2 |

```
In [221]: df_GROUPED_C_J.classss.value_counts(normalize=True)
```

```
Out[221]:
```

| | |
|---|----------|
| 3 | 0.435373 |
| 2 | 0.320101 |
| 1 | 0.244526 |

Name: classss, dtype: float64

```
In [222]: df_GROUPED_C_J.drop(columns=['count','mean','std','min','25%','50%','75%','max'],axis = 1,inplace=True)
```

```
In [223]: df_GROUPED_CTC_BY_E_P_C.iloc[0]
```

```
Out[223]:
```

| | |
|-------------------------------------|----------------|
| company_hash | atrgxnnt xzaxv |
| orgyear | 2016.0 |
| ctc | 1100000 |
| job_position | other |
| ctc_updated_year | 2020.0 |
| years_of_experience_in_organization | 7.0 |
| designation_in_organization | 2 |

Name: 0, dtype: object

```
In [ ]:
```

```
In [224]: df_GROUPED_C_J.iloc[0]
```

```
Out[224]:
```

| | |
|-------------------------------------|----------------|
| company_hash | atrgxnnt xzaxv |
| orgyear | 2016.0 |
| ctc | 1100000 |
| job_position | other |
| ctc_updated_year | 2020.0 |
| years_of_experience_in_organization | 7.0 |
| classss | 1 |

Name: 0, dtype: object

```
In [ ]:
```

```
In [225]: df_Grouped = df_GROUPED_CTC_BY_E_P_C.merge(df_GROUPED_C_J, on=['company_hash',
    'orgyear',
    'ctc',
    'job_position',
    'years_of_experience_in_organization',
    'ctc_updated_year'], how='left')
```

```
In [226]: df_Grouped.sample(5)
```

```
Out[226]:
```

| | company_hash | orgyear | ctc | job_position | ctc_updated_year | years_of_experience_in_organization | designation_in_organization | class |
|--------|----------------------------|---------|---------|--------------------|------------------|-------------------------------------|-----------------------------|-------|
| 107145 | cgjrrv evoxygzo uch rna | 2018.0 | 900000 | fullstack engineer | 2018.0 | 5.0 | 2 | 3 |
| 100926 | ztw ntwyzgrgsxto xzavx rna | 2019.0 | 540000 | Others | 2021.0 | 4.0 | 2 | 3 |
| 62274 | eoo | 2019.0 | 500000 | backend engineer | 2021.0 | 4.0 | 2 | 2 |
| 10909 | vqwtoxb | 2019.0 | 1500000 | Others | 2020.0 | 4.0 | 3 | 3 |
| 60338 | bxzanqt | 2017.0 | 488000 | support engineer | 2020.0 | 6.0 | 2 | 3 |

```
In [227]: df_Grouped.shape
```

```
Out[227]: (166228, 8)
```

Manual Clustering based on company

based on ctc per company , assigning company as tier 1 2 and 3 per each learners

```
In [228]: GROUPED_C = df.groupby(['company_hash'])['ctc'].describe()
df_company = df.merge(GROUPED_C, on=['company_hash'], how='left')
```

```
In [229]: df_company.sample(5)
```

```
Out[229]:
```

| | company_hash | orgyear | ctc | job_position | ctc_updated_year | years_of_experience_in_organization | count | mean | std | min | 25% | 50% | 75% | max |
|--------|----------------|---------|---------|-------------------|------------------|-------------------------------------|---------|--------------|--------------|----------|-----------|-----------|-----------|------------|
| 46717 | Others | 2012.0 | 700000 | ios engineer | 2019.0 | 11.0 | 26256.0 | 1.311366e+06 | 1.436286e+06 | 38000.0 | 440000.0 | 900000.0 | 1650000.0 | 12500000.0 |
| 93680 | wwwgrxnt | 2013.0 | 1900000 | backend engineer | 2018.0 | 10.0 | 165.0 | 1.414836e+06 | 6.917496e+05 | 200000.0 | 1000000.0 | 1300000.0 | 1500000.0 | 4800000.0 |
| 136812 | x vb v onhatzn | 2018.0 | 600000 | other | 2018.0 | 5.0 | 49.0 | 1.206531e+06 | 1.115764e+06 | 100000.0 | 500000.0 | 900000.0 | 1700000.0 | 6500000.0 |
| 111948 | nguuq | 2018.0 | 1600000 | Others | 2019.0 | 5.0 | 72.0 | 1.707083e+06 | 1.085002e+06 | 350000.0 | 900000.0 | 1560000.0 | 2070000.0 | 5000000.0 |
| 9856 | Others | 2016.0 | 1440000 | frontend engineer | 2019.0 | 7.0 | 26256.0 | 1.311366e+06 | 1.436286e+06 | 38000.0 | 440000.0 | 900000.0 | 1650000.0 | 12500000.0 |

```
In [230]: df_company['tier'] = df_company.apply(lambda x: classification(x['ctc'],x['50%'],x['75%']),axis=1)
```

```
In [231]: # df_company.sample(5)
```

```
In [232]: df_company.tier.value_counts(normalize=True)
```

```
Out[232]: 3    0.477364  
2    0.282911  
1    0.239725  
Name: tier, dtype: float64
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [233]: df_company.drop(['count','mean','std','min','25%','50%','75%','max'],  
axis = 1,  
inplace=True)
```

```
In [234]: df_company.iloc[0]
```

```
Out[234]: company_hash      atrgxnnnt xzaxv  
orgyear           2016.0  
ctc              1100000  
job_position     other  
ctc_updated_year 2020.0  
years_of_experience_in_organization 7.0  
tier             2  
Name: 0, dtype: object
```

```
In [235]: df_Grouped.iloc[0]
```

```
Out[235]: company_hash      atrgxnnnt xzaxv  
orgyear           2016.0  
ctc              1100000  
job_position     other  
ctc_updated_year 2020.0  
years_of_experience_in_organization 7.0  
designation_in_organization 2  
classs           1  
Name: 0, dtype: object
```

```
In [236]: df_Grouped = df_Grouped.merge(df_company,  
on=['company_hash',  
'orgyear', 'ctc',  
'job_position',  
'years_of_experience_in_organization',  
'ctc_updated_year'  
])
```

```
In [237]: df_Grouped
```

```
Out[237]:
```

| | company_hash | orgyear | ctc | job_position | ctc_updated_year | years_of_experience_in_organization | designation_in_organization | classs | tier |
|--------|-------------------------|---------|---------|--------------------|------------------|-------------------------------------|-----------------------------|--------|------|
| 0 | atrgxnt xzaxv | 2016.0 | 1100000 | other | 2020.0 | 7.0 | | 2 | 1 2 |
| 1 | qtrxzwt xzegwgb rxbxnta | 2018.0 | 449999 | fullstack engineer | 2019.0 | 5.0 | | 3 | 3 3 |
| 2 | Others | 2015.0 | 2000000 | backend engineer | 2020.0 | 8.0 | | 1 | 1 1 |
| 3 | ngpputaxv | 2017.0 | 700000 | backend engineer | 2019.0 | 6.0 | | 3 | 3 3 |
| 4 | qxen sqghu | 2017.0 | 1400000 | fullstack engineer | 2019.0 | 6.0 | | 2 | 1 1 |
| ... | ... | ... | ... | ... | ... | ... | | ... | ... |
| 171311 | vuurtxzw | 2008.0 | 220000 | Others | 2019.0 | 15.0 | | 2 | 3 3 |
| 171312 | husqyawgb | 2017.0 | 500000 | Others | 2020.0 | 6.0 | | 3 | 3 3 |
| 171313 | vwwgrxnt | 2021.0 | 700000 | Others | 2021.0 | 2.0 | | 2 | 3 3 |
| 171314 | zgn vuurxwvmt | 2019.0 | 5100000 | Others | 2019.0 | 4.0 | | 1 | 1 1 |
| 171315 | bgqsqvz onvzrtj | 2014.0 | 1240000 | Others | 2016.0 | 9.0 | | 3 | 3 3 |

171316 rows × 9 columns

```
In [238]: X = df_Grouped.copy()
```

```
In [ ]:
```

```
In [239]: X.shape
```

```
Out[239]: (171316, 9)
```

```
In [240]: X_data = X.drop(["company_hash", "job_position"], axis = 1)
```

Final data for Model :

```
In [241]: X_data
```

```
Out[241]:
```

| | orgyear | ctc | ctc_updated_year | years_of_experience_in_organization | designation_in_organization | classs | tier |
|--------|---------|---------|------------------|-------------------------------------|-----------------------------|--------|------|
| 0 | 2016.0 | 1100000 | 2020.0 | 7.0 | | 2 | 1 2 |
| 1 | 2018.0 | 449999 | 2019.0 | 5.0 | | 3 | 3 3 |
| 2 | 2015.0 | 2000000 | 2020.0 | 8.0 | | 1 | 1 1 |
| 3 | 2017.0 | 700000 | 2019.0 | 6.0 | | 3 | 3 3 |
| 4 | 2017.0 | 1400000 | 2019.0 | 6.0 | | 2 | 1 1 |
| ... | ... | ... | ... | ... | | ... | ... |
| 171311 | 2008.0 | 220000 | 2019.0 | 15.0 | | 2 | 3 3 |
| 171312 | 2017.0 | 500000 | 2020.0 | 6.0 | | 3 | 3 3 |
| 171313 | 2021.0 | 700000 | 2021.0 | 2.0 | | 2 | 3 3 |
| 171314 | 2019.0 | 5100000 | 2019.0 | 4.0 | | 1 | 1 1 |
| 171315 | 2014.0 | 1240000 | 2016.0 | 9.0 | | 3 | 3 3 |

171316 rows × 7 columns

Standardization:

```
In [242]: from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
scaler.fit(X_data)
X_sc = pd.DataFrame(scaler.transform(X_data), columns=X_data.columns, index=X_data.index)
```

```
In [ ]:
```

```
In [243]: X_sc
```

```
Out[243]:
```

| | orgyear | ctc | ctc_updated_year | years_of_experience_in_organization | designation_in_organization | classs | tier |
|--------|-----------|-----------|------------------|-------------------------------------|-----------------------------|-----------|-----------|
| 0 | 0.229439 | -0.238430 | 0.298195 | -0.229439 | -0.175910 | -1.497105 | -0.300556 |
| 1 | 0.680950 | -0.741765 | -0.452799 | -0.680950 | 1.196414 | 1.001707 | 0.933655 |
| 2 | 0.003683 | 0.458493 | 0.298195 | -0.003683 | -1.548235 | -1.497105 | -1.534766 |
| 3 | 0.455194 | -0.548174 | -0.452799 | -0.455194 | 1.196414 | 1.001707 | 0.933655 |
| 4 | 0.455194 | -0.006122 | -0.452799 | -0.455194 | -0.175910 | -1.497105 | -1.534766 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 171311 | -1.576605 | -0.919866 | -0.452799 | 1.576605 | -0.175910 | 1.001707 | 0.933655 |
| 171312 | 0.455194 | -0.703046 | 0.298195 | -0.455194 | 1.196414 | 1.001707 | 0.933655 |
| 171313 | 1.358216 | -0.548174 | 1.049190 | -1.358216 | -0.175910 | 1.001707 | 0.933655 |
| 171314 | 0.906705 | 2.859008 | -0.452799 | -0.906705 | -1.548235 | -1.497105 | -1.534766 |
| 171315 | -0.222072 | -0.130020 | -2.705782 | 0.222072 | 1.196414 | 1.001707 | 0.933655 |

171316 rows × 7 columns

hierarchical Clustering :

trying to get a high level idea about how many clusters we can form, by taking sample of 500 learners multiple times and forming hierarchy and visualising in dendrogram.

```
In [244]: import scipy.cluster.hierarchy as sch
import matplotlib.pyplot as plt

sample = X_sc.sample(500)
Z = sch.linkage(sample, method='ward')

fig, ax1 = plt.subplots(figsize=(20, 12))
sch.dendrogram(Z, labels=sample.index, ax=ax1, color_threshold=2)
plt.xticks(rotation=90)
ax1.set_ylabel('distance')
plt.show()

import scipy.cluster.hierarchy as sch
import matplotlib.pyplot as plt


sample = X_sc.sample(500)
Z = sch.linkage(sample, method='ward')


fig, ax2 = plt.subplots(figsize=(20, 12))
sch.dendrogram(Z, labels=sample.index, ax=ax2, color_threshold=2)
plt.xticks(rotation=90)
ax2.set_ylabel('distance')
plt.show()


import scipy.cluster.hierarchy as sch
import matplotlib.pyplot as plt

sample = X_sc.sample(500)
Z = sch.linkage(sample, method='ward')

fig, ax3 = plt.subplots(figsize=(20, 12))
sch.dendrogram(Z, labels=sample.index, ax=ax3, color_threshold=2)
plt.xticks(rotation=90)
ax3.set_ylabel('distance')
plt.show()
```







Based on dendrogram , we can observe there are 3 clusters in the data based on similarity

Further checking appropriate number of clusters using Elbow Method using k-Means clustering :

In []:

KMeans

```
In [ ]: for i in range(1,10):
    from sklearn.cluster import KMeans
    k = 4
    kM = KMeans(n_clusters=k,
                 random_state=654)
    y_pred = kM.fit_predict(X_sc)
```




```
In [254]: kmeans_per_k = [KMeans(n_clusters=k, random_state=42).fit(X_sc)
for k in range(1, 10)]
inertias = [model.inertia_ for model in kmeans_per_k]
inertias
```



```
Out[254]: [1199211.999999972,
812618.2236265242,
663951.3689564556,
577020.6292578052,
517714.4060221886,
476402.90178635635,
439357.96141059144,
410144.6171733509,
383988.5907258121]
```

```
In [255]: plt.figure(figsize=(12, 8))
plt.plot(range(1, 10), inertias, "bo-")
plt.xlabel("$k$")
plt.ylabel("Inertia", fontsize=14)
plt.annotate('Elbow',
            xy=(3, inertias[2]),
            xytext=(0.55, 0.55),
            textcoords='figure fraction',
            fontsize=16,
            arrowprops=dict(facecolor='black', shrink=0.1)
)
plt.show()
```



KMeans with n_clusters = 3

```
In [256]: from sklearn.cluster import KMeans
k = 3
kM = KMeans(n_clusters=k,
             random_state=654)
y_pred = kM.fit_predict(X_sc)
```

```
In [257]: clusters = pd.DataFrame(X, columns=X.columns)
clusters['label'] = kM.labels_
```

```
In [258]: clusters.sample(5)
```

```
Out[258]:
```

| | company_hash | orgyear | ctc | job_position | ctc_updated_year | years_of_experience_in_organization | designation_in_organization | classs | tier | label |
|--------|----------------------------|---------|---------|--------------------|------------------|-------------------------------------|-----------------------------|--------|------|-------|
| 69989 | Others | 2020.0 | 360000 | support engineer | 2020.0 | 3.0 | 3 | 3 | 3 | 2 |
| 160236 | otvqo ygraxzso wgqugqvnxgz | 2017.0 | 8000000 | data scientist | 2019.0 | 6.0 | 1 | 1 | 1 | 1 |
| 101242 | mvqwrjo | 2001.0 | 3350000 | Others | 2019.0 | 22.0 | 1 | 1 | 1 | 0 |
| 136293 | nvvn wzgzhmrvwj otqcxwto | 2015.0 | 1220000 | fullstack engineer | 2021.0 | 8.0 | 1 | 1 | 1 | 1 |
| 27089 | wbt sqghu | 2011.0 | 1600000 | Others | 2019.0 | 12.0 | 2 | 3 | 3 | 2 |

```
In [259]: clusters.shape
```


```
Out[259]: (171316, 10)
```

Insights | EDA after Clustering :

```
In [260]: sns.scatterplot(clusters["orgyear"],  
                      clusters["ctc"],  
                      hue = clusters["label"])
```

```
<IPython.core.display.Javascript object>
```

```
Out[260]: <AxesSubplot:xlabel='orgyear', ylabel='ctc'>
```



```
In [275]:
```

```
Out[275]: 2000000.0
```


based on above scatter plot , we can observe , a cluster of learners received CTC upto 30 LPA who joined after 2006-07.

there's a group of learners who are very much experienced.

and also learners joined after 2012-13 receiving CTC between 20 LPA to upto 1.5cr.

```
In [261]: pd.crosstab(index = clusters["label"],  
                   columns = clusters["tier"],values=clusters["ctc"],aggfunc= np.mean  
                   ).plot(kind = "bar")
```


```
Out[261]: <AxesSubplot:xlabel='label'>
```



Based on k-Means Clustering algorithm output , as well as manual clustering , learners from tier1 company receiving very high CTC.


```
In [262]: pd.crosstab(index = clusters["label"],
                   columns = clusters["classs"],values=clusters["ctc"],aggfunc= np.mean
                   ).plot(kind = "bar")
```

```
Out[262]: <AxesSubplot:xlabel='label'>
```




```
In [263]: pd.crosstab(index = clusters["label"],
                   columns = clusters["designation_in_organization"],
                   values=clusters["ctc"],aggfunc= np.mean
                  ).plot(kind = "bar")
```

```
Out[263]: <AxesSubplot:xlabel='label'>
```



```
In [264]: pd.crosstab(columns = clusters["label"],
                     index = clusters["years_of_experience_in_organization"],
                     values=clusters["ctc"],aggfunc= np.mean
                    ).plot(kind = "bar")
```


```
Out[264]: <AxesSubplot:xlabel='years_of_experience_in_organization'>
```



**Cluster label 0 , are those learners who are very very experienced,
experienced learners between 6 to 10 years of experience, earning above 40 LPA up tp 1.5Cr.**

```
In [266]: pd.crosstab(columns = clusters["label"],  
                   index = clusters["years_of_experience_in_organization"],  
                   ).plot(kind = "bar")
```

```
Out[266]: <AxesSubplot:xlabel='years_of_experience_in_organization'>
```



Majority of Learners are experienced between 1 to 15 years . (49.73%)- (Cluster 2)

there is a group of learners having 8 to upto 33 years of experience. (33%) - (Cluster 0)

16.95% of learners who have experiences - (cluster 1)

```
In [278]: clusters.label.value_counts(normalize=True)*100
```

```
Out[278]: 2    49.734409  
1    33.308623  
0    16.956968  
Name: label, dtype: float64
```


```
In [ ]:
```

```
In [ ]:
```

years_of_experience_in_organization per each cluster group of learners

```
In [269]: pd.crosstab(index = clusters["label"],
                  columns = clusters["tier"],
                  values=clusters["years_of_experience_in_organization"],
                  aggfunc=np.mean
                 ).plot(kind = "bar")
```

```
Out[269]: <AxesSubplot:xlabel='label'>
```



```
In [286]: clusters.columns
```

```
Out[286]: Index(['company_hash', 'orgyear', 'ctc', 'job_position', 'ctc_updated_year',
                  'years_of_experience_in_organization', 'designation_in_organization',
                  'classs', 'tier', 'label'],
                 dtype='object')
```

Statistical Summary based on Each Cluster :

```
In [288]: clusters.groupby("label").describe()[["ctc","classs","tier","years_of_experience_in_organization"]].T
```

```
Out[288]:
```

| | label | 0 | 1 | 2 |
|-------------------------------------|-------|--------------|--------------|--------------|
| ctc | count | 2.905000e+04 | 5.706300e+04 | 8.520300e+04 |
| | mean | 2.543348e+06 | 1.802940e+06 | 7.562107e+05 |
| | std | 1.751976e+06 | 1.272597e+06 | 5.033019e+05 |
| | min | 3.955000e+04 | 6.500000e+04 | 3.800000e+04 |
| | 25% | 1.420000e+06 | 1.000000e+06 | 4.000000e+05 |
| | 50% | 2.100000e+06 | 1.500000e+06 | 6.300000e+05 |
| | 75% | 3.147500e+06 | 2.200000e+06 | 1.000000e+06 |
| | max | 1.250000e+07 | 1.250000e+07 | 5.600000e+06 |
| classs | count | 2.905000e+04 | 5.706300e+04 | 8.520300e+04 |
| | mean | 1.625886e+00 | 1.544574e+00 | 2.831191e+00 |
| | std | 6.937293e-01 | 5.252113e-01 | 3.751798e-01 |
| | min | 1.000000e+00 | 1.000000e+00 | 1.000000e+00 |
| | 25% | 1.000000e+00 | 1.000000e+00 | 3.000000e+00 |
| | 50% | 2.000000e+00 | 2.000000e+00 | 3.000000e+00 |
| | 75% | 2.000000e+00 | 2.000000e+00 | 3.000000e+00 |
| | max | 3.000000e+00 | 3.000000e+00 | 3.000000e+00 |
| tier | count | 2.905000e+04 | 5.706300e+04 | 8.520300e+04 |
| | mean | 1.484200e+00 | 1.648774e+00 | 2.900731e+00 |
| | std | 6.478262e-01 | 5.742163e-01 | 3.010974e-01 |
| | min | 1.000000e+00 | 1.000000e+00 | 1.000000e+00 |
| | 25% | 1.000000e+00 | 1.000000e+00 | 3.000000e+00 |
| | 50% | 1.000000e+00 | 2.000000e+00 | 3.000000e+00 |
| | 75% | 2.000000e+00 | 2.000000e+00 | 3.000000e+00 |
| | max | 3.000000e+00 | 3.000000e+00 | 3.000000e+00 |
| years_of_experience_in_organization | count | 2.905000e+04 | 5.706300e+04 | 8.520300e+04 |
| | mean | 1.520678e+01 | 6.557945e+00 | 6.541436e+00 |
| | std | 4.339403e+00 | 2.474935e+00 | 2.775220e+00 |
| | min | 6.000000e+00 | 1.000000e+00 | 1.000000e+00 |
| | 25% | 1.200000e+01 | 5.000000e+00 | 4.000000e+00 |
| | 50% | 1.400000e+01 | 7.000000e+00 | 6.000000e+00 |
| | 75% | 1.700000e+01 | 8.000000e+00 | 8.000000e+00 |
| | max | 3.300000e+01 | 1.300000e+01 | 1.700000e+01 |

```
In [ ]:
```

5. Actionable Insights & Recommendations

- Maximum users have years of experience in the range of 3-10 years. Scaler can target the audience with experience 3 to 10 because mostly they were looking for a career change or upskilling.
- In data, email_hash are repeating data and should store one email_hash for the individual users.
- Maximum users have job_positions as other and backed_engineer, so scaler can target more on these 2 job_position.
- With company_hash we can see the high-paying companies for the different job roles so the scaler can target the audience or advertise with these data insights.
- In orgyear most of the years are invalid so the system should give an error if the user enters the wrong orgyear.
- Scaler should show insights about the updated ctc and use this as a marketing strategy to bring in more audience.